CS 4287
Team 10
Zhitao Shu, Duke Jiang

How to run this assignment:
1. Start VM1 on the local machine, download the repo by running sudo git clone
   https://github.com/Shuz98/CloudComputing_assignment1
2. Cd to the directory of the repo to call the producer later.
3. On the cloud, ssh into VM2, and run zookeeper server by running sudo systemctl start
   zookeeper
4. On VM2, run a Kafka broker by running sudo kafka-server-start.sh /etc/kafka.properties
5. On the cloud, ssh into VM3, and run zookeeper server by running sudo systemctl start
   zookeeper
6. On VM3, run a second Kafka broker by running sudo kafka-server-start.sh
   /etc/kafka.properties
7. Pull the githubt repo to VM3 by running sudo git clone
   https://github.com/Shuz98/CloudComputing_assignment1
8. Cd into the repo directory, run the consumer program by running python3 consumer.py
9. On the local VM1, run the producer program by running python3 producer.py